

✓

Medical Insurance Price Prediction

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib as pt
5 import warnings
6 warnings.filterwarnings("ignore")
```

```
1 df=pd.read_csv("/content/insurance.csv")
2 df
```

↗

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

1338 rows × 7 columns

```
1 df.info()
```

↗

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   expenses    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
1 df.describe()
```

↗

	age	bmi	children	expenses
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.665471	1.094918	13270.422414
std	14.049960	6.098382	1.205493	12110.011240
min	18.000000	16.000000	0.000000	1121.870000
25%	27.000000	26.300000	0.000000	4740.287500
50%	39.000000	30.400000	1.000000	9382.030000
75%	51.000000	34.700000	2.000000	16639.915000
max	64.000000	53.100000	5.000000	63770.430000

```
1 df.isnull().sum()
```

```

age      0
sex      0
bmi      0
children 0
smoker   0
region   0
expenses 0
dtype: int64


```

✓ Pie chart for the sex, smoker, and region column

```

1 features = ['sex', 'smoker', 'region']
2
3 plt.subplots(figsize=(20, 10))
4 for i, col in enumerate(features = ['age', 'bmi'])
5
6 plt.subplots(figsize=(17, 7))
7 for i, col in enumerate(features):
8     plt.subplot(1, 2, i + 1)
9     sb.scatterplot(data=df, x=col,
10                    y='charges',
11                    hue='smoker')
12 plt.show()
13 features):
14     plt.subplot(1, 3, i + 1)
15
16     x = df[col].value_counts()
17     plt.pie(x.values,
18            labels=x.index,
19            autopct='%1.1f%%')
20
21 plt.show()

```





Comparison between expenses paid between different groups

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 # Convert 'charges' column to numeric, handling errors
5 # Check if 'expenses' column exists before conversion
6 if 'expenses' in df.columns:
7     df['expenses'] = pd.to_numeric(df['expenses'], errors='coerce')
8
9 features = ['sex', 'children', 'smoker', 'region']
10
11 plt.subplots(figsize=(20, 10))
12 for i, col in enumerate(features):
13     plt.subplot(2, 2, i + 1)
14     # Handle potential missing values in 'expenses'
15     df.groupby(col)['expenses'].mean().plot.bar()
16 plt.show()


```



Scatter plot of the charges paid v/s age and BMI respectively

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sb # Import seaborn
4
5
6 df = pd.read_csv('/content/insurance.csv')
7
8 features = ['age', 'bmi']
9
10 plt.subplots(figsize=(17, 7))
11 for i, col in enumerate(features):
12     plt.subplot(1, 2, i + 1)
13     sb.scatterplot(data=df, x=col,
14                   y='expenses',
15                   hue='smoker')
16 plt.show()
```


<ipython-input-10-a85cdb1d55bd>:12: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated
plt.subplot(1, 2, i + 1)



Boxplot of age

```
1 df.drop_duplicates(inplace=True)
2 sb.boxplot(df['age'])
```


↳ <Axes: ylabel='age'>



Box plot of bmi

```
1 import seaborn as sns
2 sns.boxplot(df['bmi'])
```

↳ <Axes: ylabel='bmi'>



```
1 Q1=df['bmi'].quantile(0.25)
2 Q2=df['bmi'].quantile(0.5)
3 Q3=df['bmi'].quantile(0.75)
4 iqr=Q3-Q1
5 lowlim=Q1-1.5*iqr
6 upplim=Q3+1.5*iqr
7 print(lowlim)
8 print(upplim)
```

↳ 13.699999999999998
47.300000000000004

```
1 df['age'].skew()
```


↳ 0.054780773126998195

correlation matrix

```

1 !pip install --upgrade pandas
2
3 import pandas as pd
4 import matplotlib.pyplot as plt
5
6
7
8 df['sex'] = df['sex'].map({'male': 0, 'female': 1})
9 df['smoker'] = df['smoker'].map({'yes': 1, 'no': 0})
10 df['region'] = df['region'].map({'northwest': 0, 'northeast': 1, 'southeast': 2, 'southwest': 3})
11
12 # Try using to_string() to display the correlation matrix
13 print(df.corr().to_string())

```


 Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)
 Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.25.2)
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas)

	age	sex	bmi	children	smoker	region	expenses
age	1.000000	NaN	0.112069	0.041536	NaN	NaN	0.298308
sex	NaN	NaN	NaN	NaN	NaN	NaN	NaN
bmi	0.112069	NaN	1.000000	0.013574	NaN	NaN	0.199298
children	0.041536	NaN	0.013574	1.000000	NaN	NaN	0.067389
smoker	NaN	NaN	NaN	NaN	NaN	NaN	NaN
region	NaN	NaN	NaN	NaN	NaN	NaN	NaN
expenses	0.298308	NaN	0.199298	0.067389	NaN	NaN	1.000000

```

1 xtrain,xtest,ytrain,ytest=train_test_split(X,Y,test_size=0.2,random_state=42)
2 lrmodel=LinearRegression()
3 lrmodel.fit(xtrain,ytrain)
4 print(lrmodel.score(xtrain,ytrain))
5 print(lrmodel.score(xtest,ytest))
6 print(cross_val_score(lrmodel,X,Y,cv=5,).mean())


```

 0.1096025209475574
 0.13771209251047456
 0.1167572302298773

```

1 from sklearn.metrics import r2_score
2 svrmodel=SVR()
3 svrmodel.fit(xtrain,ytrain)
4 ypredtrain1=svrmodel.predict(xtrain)
5 ypredtest1=svrmodel.predict(xtest)
6 print(r2_score(ytrain,ypredtrain1))
7 print(r2_score(ytest,ypredtest1))
8 print(cross_val_score(svrmodel,X,Y,cv=5,).mean())

```

 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)
 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)
 -0.1006002667676189
 -0.13368071493013267
 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)
 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)
 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)
 -0.10361331629076327
 /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1310: DataConversionWarning: A column-vector y was passed as a 1D array, which will be converted to a column vector by default.
 y = column_or_1d(y, warn=True)

```
1 rfmodel=RandomForestRegressor(random_state=42)
2 rfmodel.fit(xtrain,ytrain)
3 ypredtrain2=rfmodel.predict(xtrain)
4 ypredtest2=rfmodel.predict(xtest)
5 print(r2_score(ytrain,ypredtrain2))
6 print(r2_score(ytest,ypredtest2))
7 print(cross_val_score(rfmodel,X,Y,cv=5,).mean())
8 from sklearn.model_selection import GridSearchCV
9 estimator=RandomForestRegressor(random_state=42)
10 param_grid={'n_estimators':[10,40,50,98,100,120,150]}
11 grid=GridSearchCV(estimator,param_grid,scoring="r2",cv=5)
12 grid.fit(xtrain,ytrain)
13 print(grid.best_params_)
14 rfmodel=RandomForestRegressor(random_state=42,n_estimators=120)
15 rfmodel.fit(xtrain,ytrain)
16 ypredtrain2=rfmodel.predict(xtrain)
17 ypredtest2=rfmodel.predict(xtest)
18 print(r2_score(ytrain,ypredtrain2))
19 print(r2_score(ytest,ypredtest2))
20 print(cross_val_score(rfmodel,X,Y,cv=5,).mean())
```

RandomForestRegressor:

0.9738163260247533

0.8819423353068565

0.8363637309718952

Hyperparameter tuning:

{'n_estimators': 120}

0.9746383984429655

0.8822009842175969

0.8367438097052858

```
1 gbmodel=GradientBoostingRegressor()
2 gbmodel.fit(xtrain,ytrain)
3 ypredtrain3=gbmodel.predict(xtrain)
4 ypredtest3=gbmodel.predict(xtest)
5 print(r2_score(ytrain,ypredtrain3))
6 print(r2_score(ytest,ypredtest3))
7 print(cross_val_score(gbmodel,X,Y,cv=5,).mean())
8 from sklearn.model_selection import GridSearchCV
9 estimator=GradientBoostingRegressor()
10 param_grid={'n_estimators':[10,15,19,20,21,50], 'learning_rate':[0.1,0.19,0.2,0.21,0.8,1]}
11 grid=GridSearchCV(estimator,param_grid,scoring="r2",cv=5)
12 grid.fit(xtrain,ytrain)
13 print(grid.best_params_)
14 gbmodel=GradientBoostingRegressor(n_estimators=19,learning_rate=0.2)
15 gbmodel.fit(xtrain,ytrain)
16 ypredtrain3=gbmodel.predict(xtrain)
17 ypredtest3=gbmodel.predict(xtest)
18 print(r2_score(ytrain,ypredtrain3))
19 print(r2_score(ytest,ypredtest3))
20 print(cross_val_score(gbmodel,X,Y,cv=5,).mean())
```

GradientBoostingRegressor:

0.8931345821166041

0.904261922040551

0.8549940291799407

Hyperparameter tuning

{'learning_rate': 0.2, 'n_estimators': 21} 0.8682397447116927

0.9017109716082661

0.8606041910125791

```

1 xgmodel=XGBRegressor()
2 xgmodel.fit(xtrain,ytrain)
3 ypredtrain4=xgmodel.predict(xtrain)
4 ypredtest4=xgmodel.predict(xtest)
5 print(r2_score(ytrain,ypredtrain4))
6 print(r2_score(ytest,ypredtest4))
7 print(cross_val_score(xgmodel,X,Y,cv=5,).mean())
8 from sklearn.model_selection import GridSearchCV
9 estimator=XGBRegressor()
10 param_grid={'n_estimators':[10,15,20,40,50], 'max_depth':[3,4,5], 'gamma':[0,0.15,0.3,0.5,1]}
11 grid=GridSearchCV(estimator,param_grid,scoring="r2",cv=5)
12 grid.fit(xtrain,ytrain)
13 print(grid.best_params_)
14 xgmodel=XGBRegressor(n_estimators=15,max_depth=3,gamma=0)
15 xgmodel.fit(xtrain,ytrain)
16 ypredtrain4=xgmodel.predict(xtrain)
17 ypredtest4=xgmodel.predict(xtest)
18 print(r2_score(ytrain,ypredtrain4))
19 print(r2_score(ytest,ypredtest4))
20 print(cross_val_score(xgmodel,X,Y,cv=5,).mean())

```

```

0.9118184447288513
-0.2653837203979492
-0.2806396007537842
{'gamma': 0, 'max_depth': 3, 'n_estimators': 10}
0.24651789665222168
0.07248115539550781
0.053965306282043456

```

```

1 import pandas as pd # Import pandas
2
3 # Check if 'feature_importances_' exists
4 if hasattr(grid.best_estimator_, 'feature_importances_'):
5     # Handle cases where X might not have the 'columns' attribute
6     if hasattr(X, 'columns'):
7         feats = pd.DataFrame(data=grid.best_estimator_.feature_importances_,
8                               index=X.columns,
9                               columns=['Importance']) # Set column name here
10    else:
11        feats = pd.DataFrame(data=grid.best_estimator_.feature_importances_,
12                              index=range(len(grid.best_estimator_.feature_importances_)),
13                              columns=['Importance']) # Set column name here
14    print(feats) # Display the feature importances
15 else:
16    print("The 'best_estimator_' does not have the 'feature_importances_' attribute.")

```

```

Importance
0    0.525635
1    0.261984
2    0.212381

```

```

1 df.drop(df[['sex','region']],axis=1,inplace=True)
2 Xf=df.drop(df[['charges']],axis=1)
3 X=df.drop(df[['charges']],axis=1)
4 xtrain,xtest,ytrain,ytest=train_test_split(Xf,Y,test_size=0.2,random_state=42)
5 finalmodel=XGBRegressor(n_estimators=15,max_depth=3,gamma=0)
6 finalmodel.fit(xtrain,ytrain)
7 ypredtrain4=finalmodel.predict(xtrain)
8 ypredtest4=finalmodel.predict(xtest)
9 print(r2_score(ytrain,ypredtrain4))
10 print(r2_score(ytest,ypredtest4))
11 print(cross_val_score(finalmodel,X,Y,cv=5,).mean())

```

Final Model:

Train accuracy : 0.870691899927822

Test accuracy : 0.904151903449132

CV Score : 0.8600710679082143

